

DD2360

Assignment 4 - Report

Heng Kang
hengkang@kth.se

Jian Wen
jianwe@kth.se

Git repo: <https://github.com/SinGuLaRiT2001/DD2360HT23>

Ex1. Thread Scheduling and Execution Efficiency

1. Assume $X=800$ and $Y=600$. Assume that we decided to use a grid of 16×16 blocks. That is, each block is organized as a 2D 16×16 array of threads. How many warps will be generated during the execution of the kernel? How many warps will have control divergence? Please explain your answers.

ANSWER: For each block, there are $16 \times 16 = 256$ threads. The block number is $(800/16) \times (600/16) = 50 \times 38 = 1900$ blocks. Therefore, the total number of threads is $1900 \times 256 = 486400$ threads. In CUDA devices, each warp contains 32 threads, so the number of warps is $486400/32 = 15200$.

Control divergence occurs in the last row and last column of blocks due to some threads being out of the image bounds. On the X axis, the number of warp with divergence is $(800\%16)/32 = 0$. On the Y axis, the number of warp with divergence is $(600\%16)/32 = 1$. Considering there are 38 blocks on the Y axis, so there are a total of 38 warps with divergence.

2. Now assume $X=600$ and $Y=800$ instead, how many warps will have control divergence? Please explain your answers.

ANSWER: By using similar calculations, the number of warp on the X axis will be $(600\%16)/32 = 1$ while for Y axis they are perfectly aligned. And there are $600/16 = 38$ blocks on X axis. So the total number of warps is still 38.

3. Now assume $X=600$ and $Y=799$, how many warps will have control divergence? Please explain your answers.

ANSWER: By using similar calculations, the number of warp on the X axis will be $(600\%16)/32 = 1$, and the total number of warps will be $1 \times 38 = 38$ on X axis; For Y axis, there are $799/16 = 50$ blocks. And the warps number will be $50 \times (799\%16)/32 = 50 \times 8 = 400$ on Y axis. Therefore, the total number of warps with control divergence is $400 + 38 = 438$.

Ex2. CUDA Streams

Updated Implementation:


```
cudaStream_t streams[numStreams];
for (int i = 0; i < numStreams; i++) {
    cudaStreamCreate(&streams[i]);
}

startTime = getTime();
for (int i = 0; i < inputLength; i += S_seg) {
    int segSize = min(S_seg, inputLength - i);
    int streamIdx = i / S_seg % numStreams;
    cudaMemcpyAsync(deviceInput1 + i, hostInput1 + i, segSize * sizeof(DataType),
        cudaMemcpyHostToDevice, streams[streamIdx]);
    cudaMemcpyAsync(deviceInput2 + i, hostInput2 + i, segSize * sizeof(DataType),
        cudaMemcpyHostToDevice, streams[streamIdx]);
    int Block_dim = ThreadNum;
    int Grid_dim = segSize / ThreadNum + (segSize % ThreadNum != 0);
    vecAdd<<<Grid_dim, Block_dim, 0, streams[streamIdx]>>>(deviceInput1 + i, deviceInput2 + i,
        deviceOutput + i, segSize);
    cudaMemcpyAsync(hostOutput + i, deviceOutput + i, segSize * sizeof(DataType),
        cudaMemcpyDeviceToHost, streams[streamIdx]);
}
for (int i = 0; i < numStreams; i++) {
    cudaStreamSynchronize(streams[i]);
}

endTime = getTime();
printf("Total asynchronous operation time: %.6lf ms.\n", (endTime - startTime) * 1000);
```

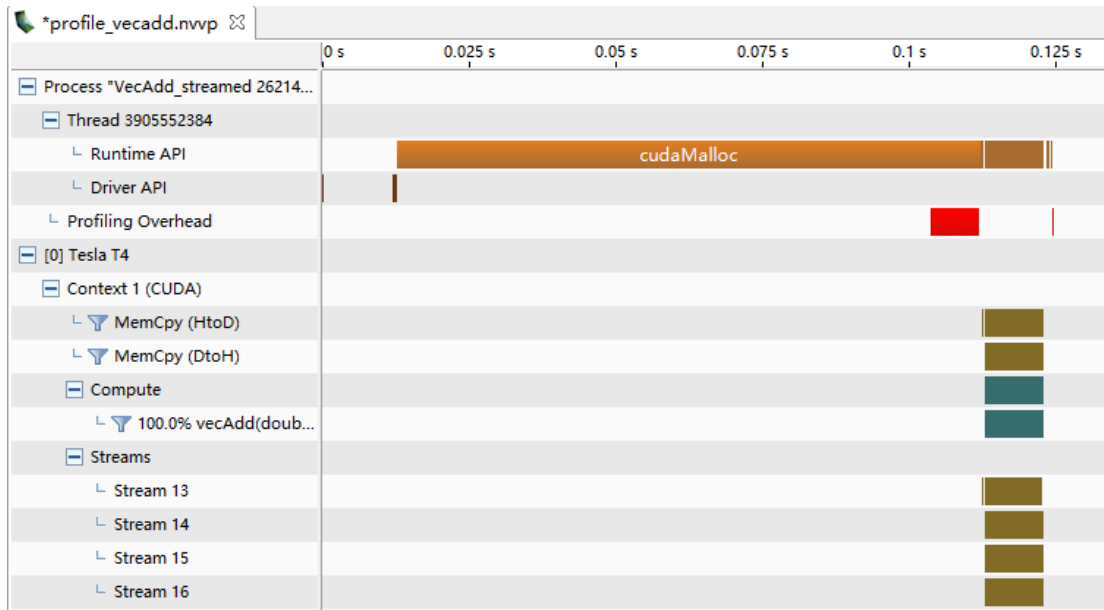
1. Compared to the non-streamed vector addition, what performance gain do you get? Present in a plot (you may include comparison at different vector length)

ANSWER: The two implementations are evaluated with different input lengths: 4096, 16384, 65536, 262144, 1048576. The S_seg is set as 1024.




- Use nvprof to collect traces and the NVIDIA Visual Profiler (nvvp) to visualize the overlap of communication and computation.

ANSWER: The NVVP timeline indicates the use of four CUDA streams. It can be seen that the memory copy and computation are carried out at the same time on 4 different streams.



- What is the impact of segment size on performance? Present in a plot (you may choose a large vector and compare 4-8 different segment sizes)

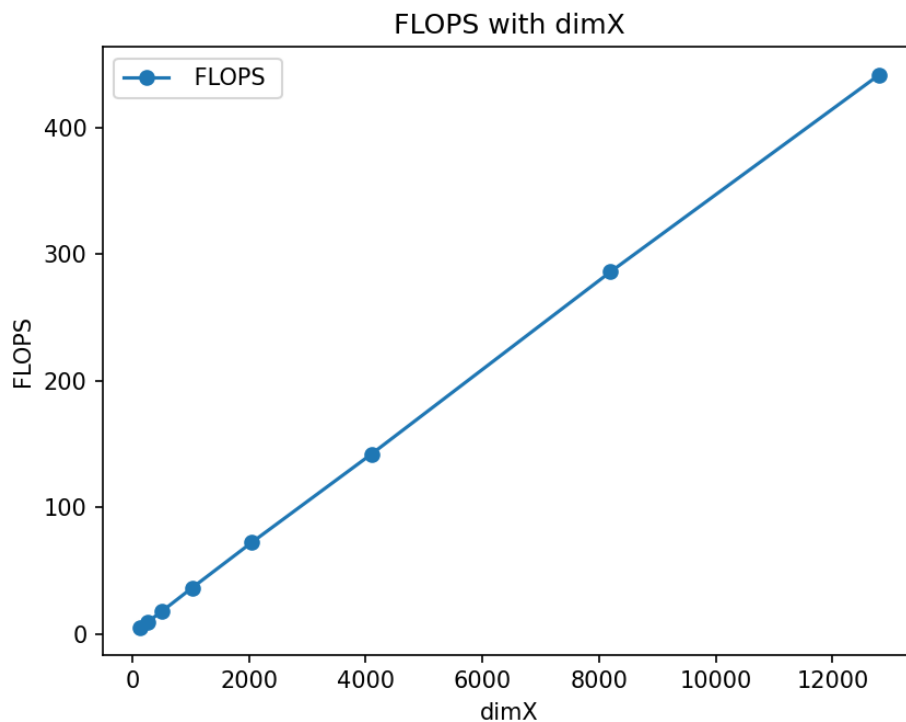
ANSWER: A series evaluation with $S_{seg}=1024, 2048, 4096, 8192, 16384, 32768, 65536$ are performed. It can be seen from the figure that, the execution time decreases as segment size increases. The x-axis of the figure is in log unit.



Ex3. Heat Equation with Using NVIDIA Libraries

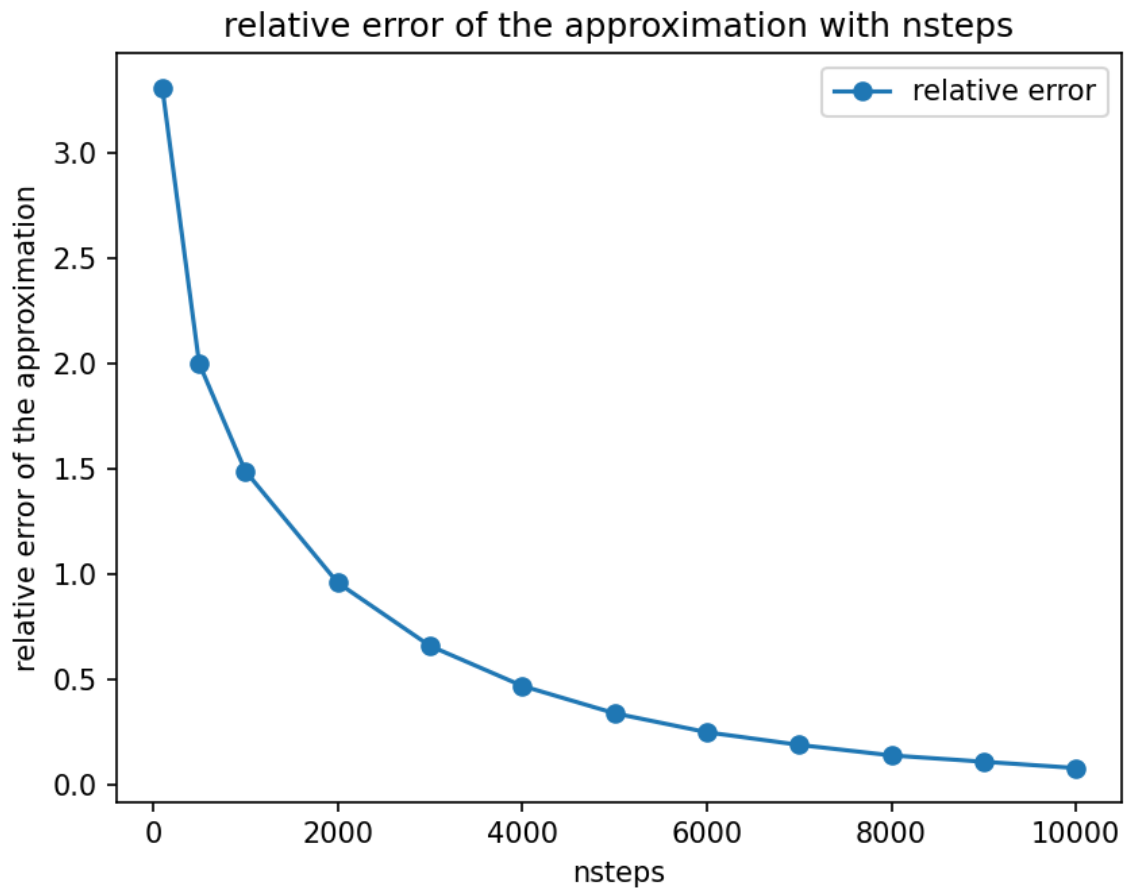
1. Run the program with different dimX values. For each one, approximate the FLOPS (floating-point operation per second) achieved in computing the SMPV (sparse matrix multiplication). Report FLOPS at different input sizes in a FLOPS. What do you see compared to the peak throughput you report in Lab2?

ANSWER: Values in the figure are lower than theoretical values.



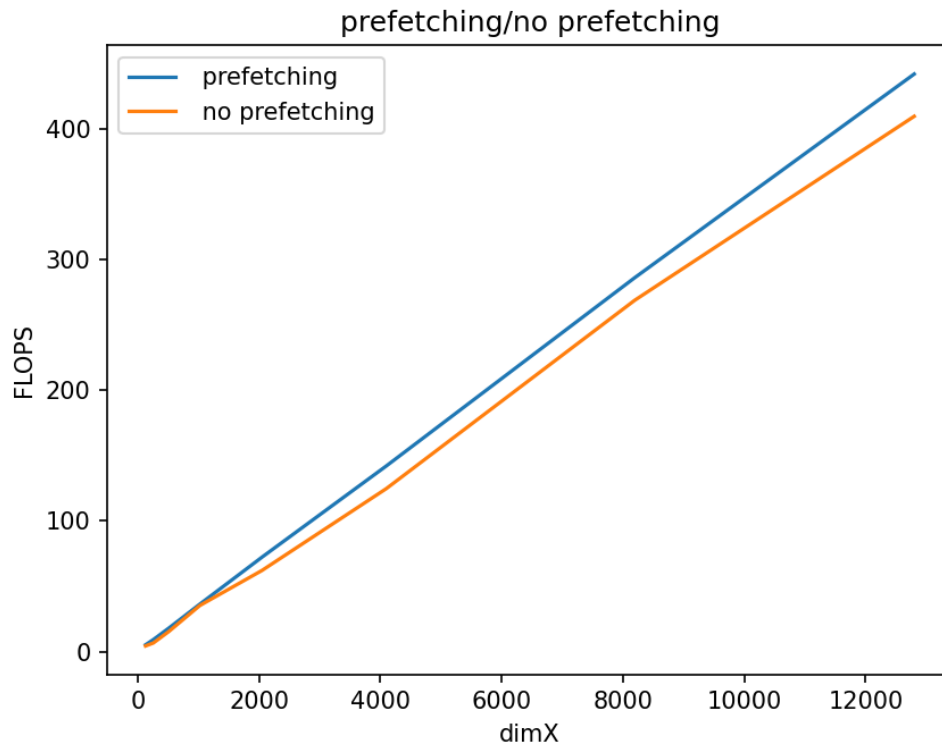
2. Run the program with dimX=128 and vary nsteps from 100 to 10000. Plot the relative error of the approximation at different nstep. What do you observe?

ANSWER: As the nstep increases, the relative error of the approximation decreases.



3. Compare the performance with and without the prefetching in Unified Memory. How is the performance impact?

ANSWER: Without prefetching, the code performs worse than prefetching.



Contributions

Jian Wen

Implemented the exercise 1&2 (Thread Scheduling and Execution Efficiency, CUDA Streams) in the assignment.

Heng Kang

Implemented the exercise 3 (Heat Equation with using NVIDIA libraries) in the assignment.